

Assignment Three

Ethan Morton

Ethan.Morton1@Marist.edu

November 15, 2024

1 GRAPHS

Template specializations are used to implement three types of undirected, non-weighted graphs: matrix, adjacency list, and linked objects. All graphs use some data structure to keep track of the vertex IDs and the edges between them. All graphs also use a lookup-table which allows you to look up the vertex in that data structure in constant time when given the vertex ID. The lookup-table is defined as `unordered_map<string, int> id_lookup;` for each implementation. The `int` represents the index used to find the vertex in a vector. This index auto-increments when a new vertex ID is added. This means the order of vertices in the vector is the same as which they appear in the graph file.

1.1 MATRIX

The data structure for the matrix is `vector<pair<string, vector<bool>>>`. To break it down, the outer vector represents each row in the matrix. The pair holds the vertex ID in the string, and the `vector<bool>` holds whether there is an edge to the vertex with that index in the vector. Initially, I was going to use `vector<vector<uint8_t>>` and do bit manipulation myself to make the matrix memory efficient, but C++ does this optimization by default when using `vector<bool>`.

Since this is a non-weighted graph, each edge only needs to store whether it exists, hence the `bool`, as opposed to a numeric type. Additionally, since the graph is undirected, it only needs to store half of the edges.

The edge between 1 and 2 is the same as the edge between 2 and 1, so we only need to track this one time. To make this optimization, the number of columns in each row is equal to that row number + 1. We have to include the +1 since a vertex can have an edge with itself. This gives us a triangle-matrix instead of a square-matrix.

When doing this optimization, we also need to consider what happens when we add an edge and the vertex IDs are given in an order we don't want. Since half the matrix is missing, we would be indexing the vector outside of its range, which is undefined behavior (or maybe C++ throws an exception... I can't remember). We can simply get around this by forcing the row to be greater than or equal to the column, and swapping them if this is false.

1.1.1 ADD VERTEX

graph.hpp

```
45 void add_vertex(const string& id) {
46     const int index = static_cast<int>(id_lookup.size());
47     id_lookup[id] = index;
48     pair<string, vector<bool>> p(id, vector<bool>(matrix.size() + 1));
49     matrix.emplace_back(p);
50 }
```

Adding a vertex is constant time. It auto-increments the index based on the size of the lookup table, then inserts the ID-index pair into the table, then pushes the ID-row pair into the matrix.

1.1.2 ADD EDGE

graph.hpp

```
52 void add_edge(const string& id0, const string& id1) {
53     uint8_t r = id_lookup[id0];
54     uint8_t c = id_lookup[id1];
55     //In an undirected graph, we only need to keep track of the edge in
56     → one spot.
57     //This is done by forcing the ordering of the given vertex IDs to
58     → always be a specific way.
59     //In this case, the row must not be less than the column.
60     //This is why we can use the triangle shape shown in the
61     → Graph<Matrix> documentation.
62     if (r < c) {
63         std::swap(r, c);
64     }
65     matrix[r].second[c] = true;
66 }
```

Adding an edge is constant time. It gets the index from the lookup-table, forces the ordering of the index, then sets the boolean to true.

1.1.3 PRINT

graph.hpp

```
65 void print() const {
66     printf(" ");
67     //print X-axis vertex IDs
68     for (const pair<string, vector<bool>>& pair : matrix) {
69         if (pair.first.length() == 1) {
70             printf(" %s", pair.first.c_str());
71         } else {
72             printf(" %s", pair.first.c_str());
73         }
74     }
75 }
```

```

73     }
74 }
75 printf("\n");
76 //Print each row
77 for (const pair<string, vector<bool>>& pair : matrix) {
78     //print vertex ID
79     if (pair.first.length() == 1) {
80         printf(" %s ", pair.first.c_str());
81     } else {
82         printf("%s ", pair.first.c_str());
83     }
84     //Print edges
85     for (const bool col : pair.second) {
86         if (col) {
87             printf("1 ");
88         } else {
89             printf(". ");
90         }
91     }
92     printf("\n");
93 }
94 }

```

The time complexity for printing the matrix is $O(\frac{n(n+1)}{2})$ or $O(n^2)$. Rows are printed in the order they appear in the graphs file. It should come out to look like a triangle.

1.2 ADJACENCY LIST

The adjacency list specialization is actually very similar to the matrix specialization. However, instead of storing whether there is an edge between each pair of vertices, it only stores the index of the vertices that each vertex has an edge with. The data structure for the adjacency list is `vector<pair<string, vector<int>>>`. The difference between this and the matrix is the `int` instead of `bool`.

Like the matrix, the outer vector represents the rows. The string is the vertex ID, and the `vector<int>` is a list of the index of each vertex that this one has an edge with.

Also like the matrix, the adjacency list uses a memory-optimization thanks to the graph being undirected. We only need to store each edge once. When adding an edge, we can look at the indices given in the function arguments. We can make sure the first one is greater than or equal to the second one, and swap them if this is false. When we update the table, we are always indexing by the larger index, and inserting the smaller index into that vertex's adjacency list. Vertex 2 will reference vertex 1 as a neighbor, but not the other way around.

If we wanted to search whether vertex 1 is adjacent to vertex 2, this memory-optimization would slightly slow down the search time, but it would still remain $O(n)$. Fortunately we aren't doing this search, so we can do this optimization cost-free.

1.2.1 ADD VERTEX

graph.hpp

```
115 void add_vertex(const string& id) {
116     id_lookup[id] = static_cast<int>(id_lookup.size());
117     pair<string, vector<int>> p(id, vector<int>());
118     adj_list.emplace_back(p);
119 }
```

Similar to the matrix specialization, adding a vertex is constant time, and it just inserts the ID-index pair into the table, and pushes the ID-list pair into the list.

1.2.2 ADD EDGE

graph.hpp

```
121 void add_edge(const string& id0, const string& id1) {
122     int index0 = id_lookup[id0];
123     int index1 = id_lookup[id1];
124     //In an undirected graph, we only need to keep track of the edge in one
125     → spot.
126     //This is done by forcing the ordering of the given vertex IDs to always
127     → be a specific way.
128     //In this case, index0 must not be less than index1.
129     //This way, the adjacency list is always indexed by the same number when
130     → given the same two vertices.
131     if (index0 < index1) {
132         std::swap(index0, index1);
133     }
134     adj_list[index0].second.emplace_back(index1);
135 }
```

Adding an edge is done in constant time. It forces the ordering of the index if necessary, then uses the larger index to index into the outer vector, then pushes the smaller index to the back of the inner vector.

1.2.3 PRINT

graph.hpp

```
134 void print() const {
135     for (const pair<string, vector<int>>& row : adj_list) {
136         //print vertex ID
137         printf("%s |", row.first.c_str());
138         //print vertex ID of adjacent vertices
139         for (const int col : row.second) {
140             printf(" %s,", adj_list[col].first.c_str());
141         }
142         printf("\n");
143     }
144 }
```

```

143     }
144 }

```

1.3 LINKED OBJECTS

The linked objects specialization is a bit different from the other two. This one keeps a vector of vertices. The Vertex struct has string ID and vector<const Vertex*> neighbors. The pointers in neighbors are pointers to the other vertices in the graph's vector.

1.3.1 ADD VERTEX

graph.hpp

```

159 void add_vertex(const string& id) {
160     id_lookup[id] = static_cast<int>(id_lookup.size());
161     Vertex v = {id, vector<const Vertex*>()};
162     vertices.emplace_back(v);
163 }

```

Similar to the other specializations, adding a vertex is done in constant time. The lookup table is updated and the vertex is added to the vertex vector.

1.3.2 ADD EDGE

graph.hpp

```

165 void add_edge(const string& id0, const string& id1) {
166     Vertex* v0 = &vertices[id_lookup[id0]];
167     Vertex* v1 = &vertices[id_lookup[id1]];
168     v0->neighbors.emplace_back(v1);
169     v1->neighbors.emplace_back(v0);
170 }

```

Adding an edge is done in constant time. Unlike the other two specializations, this one requires both vertices to be updated. The IDs are used to get the indices in the lookup table, and the addresses of those vertices are pushed into the other vertex's neighbors vector.

1.3.3 BREADTH-FIRST TRAVERSAL

graph.hpp

```

172 void print_breadth_first() const {
173     if (vertices.empty()) {return;}
174     unordered_set<const Vertex*> processed;
175     queue<const Vertex*> q;
176     const Vertex* v0 = &vertices[0];
177     processed.insert(v0);
178     q.push(v0);
179     while (!q.empty()) {
180         const Vertex* curr = q.front();

```

```

181         q.pop();
182         printf("%s, ", curr->id.c_str());
183         for (const Vertex* n : curr->neighbors) {
184             if (!processed.contains(n)) {
185                 q.push(n);
186                 processed.insert(n);
187             }
188         }
189     }
190     printf("\n");
191 }

```

Breadth-First traversal has a time complexity of $O(V + E)$ because it has to visit every vertex, and it needs to check every edge to see if a vertex has been processed. The function uses an `unordered_set` to keep track of all the pointers that have been processed. The algorithm starts by printing the vertex at index 0 in the vector. It uses a queue to keep track of which "layer" of vertices it's on. Once the first vertex is printed, it runs through all of its neighbors and prints those. The neighbors are stored in the queue, so in the next iteration, it goes through the items in the queue and prints it. A queue is used to print out "layers" at a time, where a layer represents the neighbors of the previous layer.

1.3.4 DEPTH-FIRST TRAVERSAL

graph.hpp

```

193 void print_depth_first() const {
194     unordered_set<const Vertex*> processed;
195     dft(&vertices[0], &processed);
196     printf("\n");
197 }
198
199 private:
200 void dft(const Vertex* v, unordered_set<const Vertex*>& processed) const {
201     if (!processed->contains(v)) {
202         printf("%s, ", v->id.c_str());
203         processed->insert(v);
204     }
205     for (const Vertex* n : v->neighbors) {
206         if (!processed->contains(n)) {
207             dft(n, processed);
208         }
209     }
210 }

```

Depth-First traversal also has a time complexity of $O(V + E)$ for the same reason as breadth-first traversal. Like BFT, DFT uses an `unordered_set` to track the processed vertex pointers. BFT prints all the neighbors of a vertex before moving on. DFT uses the call stack to move from vertex to vertex with each call. It recursively prints the neighbor of the vertex, then its neighbor, then its neighbor, etc. It processes the next neighbor of a vertex once the function returns that printed the previous neighbor.

2 BINARY SEARCH TREE

This is just a standard binary search tree with functions to insert, search, and perform a depth-first traversal.

2.1 INSERTION

binary_search_tree.cpp

```
24 void BinarySearchTree::insert(const string data) {
25     printf("Inserting %s - Root", data.c_str());
26     if (root == nullptr) {
27         root = new Node(data);
28         printf("\n");
29         return;
30     }
31     Node* curr = root;
32     while (curr != nullptr) {
33         if (data < curr->data) {
34             printf(",L");
35             if (curr->left == nullptr) {
36                 curr->left = new Node(data);
37                 break;
38             }
39             curr = curr->left;
40         } else {
41             printf(",R");
42             if (curr->right == nullptr) {
43                 curr->right = new Node(data);
44                 break;
45             }
46             curr = curr->right;
47         }
48     }
49     printf("\n");
50 }
```

Insertion is a $O(\log_2(n))$ operation because it has to do a binary search through the tree to find the location that the new node will be placed. It's actually slightly more because the tree likely won't be balanced. During the search, L or R is printed out to show the path of the search.

2.2 DEPTH-FIRST TRAVERSAL

binary_search_tree.cpp

```
52 void traverse_helper(const Node* curr) {
53     if (curr == nullptr) {return;}
54     traverse_helper(curr->left);
```

```

55     printf(" - %s", curr->data.c_str());
56     traverse_helper(curr->right);
57 }
58
59 void BinarySearchTree::depth_first_traverse() const {
60     printf("Depth-first traversal:");
61     traverse_helper(root);
62     printf("\n");
63 }

```

Depth-first traversal uses recursion to print out each element in the BST in order. It starts by making recursive calls to the left node, then when the function calls return, it starts printing out the nodes. Then it makes recursive calls to the right. This is how it prints from left to right.

2.3 SEARCH

binary_search_tree.cpp

```

65 int BinarySearchTree::search(const string target) const {
66     printf("Searching %s - Root", target.c_str());
67     int comparisons = 0;
68     const Node* curr = root;
69     while (curr != nullptr) {
70         comparisons++;
71         if (target == curr->data) {
72             return comparisons;
73         } else if (target < curr->data) {
74             printf(",L");
75             curr = curr->left;
76         } else {
77             printf(",R");
78             curr = curr->right;
79         }
80     }
81     return -1;
82 }

```

This is a binary search function for the binary search tree. It has a time complexity of $O(\log_2(n))$. The average number of comparisons observed for searching for the magicitems-find-in-bst.txt items is 11.952381. $\log_2(666) = 9.379378$. The reason for this discrepancy is that the tree is unbalanced, and many searches require more than 10 comparisons. The way it searches is by iterating down a specific branch until it finds the target. If it hasn't found the target, it will move to the left if the target is less than the current node, or it will move to the right if the target is greater than or equal to the current node.

3 PUTTING IT ALL TOGETHER

3.1 VERIFYING GRAPH FILES

main.cpp

```
30 int main(int argc, char* argv[]) {
31     //Make a list of all the graph files to use
32     vector<string> graph_files;
33     if (argc == 1) {
34         graph_files.emplace_back("./graphs1.txt");
35     } else {
36         for (int i = 1; i < argc; i++) {
37             struct stat buffer;
38             // File exists and is a regular file
39             if (stat(argv[i], &buffer) == 0 && (buffer.st_mode & S_IFREG) !=
40                 → 0) {
41                 graph_files.emplace_back(argv[i]);
42             } else {
43                 fprintf(stderr, "Error - invalid file path: %s\n", argv[i]);
44                 return 1;
45             }
46         }
47     }
48     //...
49 }
```

To make it easier to run the program with different graph files, I have made it so you can list the files in the arguments. See README.md for details.

3.2 GRAPHS

main.cpp

```
45 int main(int argc, char* argv[]) {
46     //...
47
48     //Do graph stuff for each graph file
49     string line;
50     for (const string& file : graph_files) {
51         std::ifstream graph_file(file);
52         if (!graph_file.is_open()) {
53             fprintf(stderr, "Unable to open file %s\n", file.c_str());
54             return 1;
55         }
56         Graph<Matrix> graph_m;
```

```

57     Graph<AdjList> graph_al;
58     Graph<vector<Vertex>> graph_lo;
59     bool init = false;
60     int lineNum = 0;
61     while (std::getline(graph_file, line)) {
62         lineNum++;
63         //ignore comments
64         if (size_t pos = line.find("--"); pos != std::string::npos) {
65             line = line.substr(0, pos);
66         }
67         //skip empty lines
68         if (line.empty()) {continue;}
69
70         //parse tokens
71         std::istringstream iss(line);
72         std::string command;
73         iss >> command;
74         if (command == "new") {
75             std::string graph;
76             iss >> graph;
77             if (graph == "graph") {
78                 if (init) {
79                     print_graphs(&graph_m, &graph_al, &graph_lo);
80                 }
81                 graph_m = Graph<Matrix>();
82                 graph_al = Graph<AdjList>();
83                 graph_lo = Graph<vector<Vertex>>();
84                 init = true;
85             } else {
86                 fprintf(stderr, "Error - unrecognized token: %s\nLine
87                     ↪ number: %d\n", graph.c_str(), lineNum);
88                 return 1;
89             }
90         } else if (command == "add") {
91             std::string type;
92             iss >> type;
93             if (type == "vertex") {
94                 if (string vertex; iss >> vertex) {
95                     graph_m.add_vertex(vertex);
96                     graph_al.add_vertex(vertex);
97                     graph_lo.add_vertex(vertex);
98                 } else {
99                     fprintf(stderr, "Error - failed to parse vertex id
100                         ↪ as int\nLine Number: %d\n", lineNum);

```

```

99         return 1;
100     }
101     } else if (type == "edge") {
102         char dash;
103         if (string vertex0, vertex1; iss >> vertex0 >> dash >>
104             ↪ vertex1 && dash == '-') {
105             graph_m.add_edge(vertex0, vertex1);
106             graph_al.add_edge(vertex0, vertex1);
107             graph_lo.add_edge(vertex0, vertex1);
108         } else {
109             fprintf(stderr, "Error - failed to parse vertex id
110                 ↪ as int\nLine Number: %d\n", lineNumber);
111             return 1;
112         }
113     } else {
114         fprintf(stderr, "Error - unrecognized token: %s\nLine
115             ↪ Number%d\n", type.c_str(), lineNumber);
116         return 1;
117     }
118 }
119 }
120 print_graphs(&graph_m, &graph_al, &graph_lo);
121 graph_file.close();
122 }
123
124 //...
125 }

```

main.cpp

```

18 void print_graphs(const Graph<Matrix>* graph_m, const Graph<AdjList>*
19     ↪ graph_al, const Graph<vector<Vertex>>* graph_lo) {
20     printf("Graph with Matrix:\n");
21     graph_m->print();
22     printf("\nGraph with Adjacency List:\n");
23     graph_al->print();
24     printf("\nGraph with Linked Objects - Depth First:\n");
25     graph_lo->print_depth_first();
26     printf("\nGraph with Linked Objects - Breadth First:\n");
27     graph_lo->print_breadth_first();
28     printf("\n");
29 }

```

This chunk of code is responsible for going through each graph file and doing the following:

- Going line-by-line in the text file
- Parsing and tokenizing
- Handling the few possible tokens in a syntax tree

Here are the instructions handled in the syntax tree:

- "new graph" or reaching the end of the file calls the `print_graphs` function to print the 3 graphs.
- "add vertex #" calls the `add_vertex` function on the 3 graphs. # represents the vertex ID.
- "add edge # - #" calls the `add_edge` function on the 3 graphs. The #'s represent the vertex IDs.
- Unrecognized tokens will cause the program to return early with an error.

3.3 BINARY SEARCH TREE INSERTION AND TRAVERSAL

main.cpp

```
121 int main(int argc, char* argv[]) {
122     //...
123
124     //Read lines from magicitems.txt and insert into BST
125     std::ifstream magicitems("../magicitems.txt");
126     if (!magicitems.is_open()) {
127         fprintf(stderr, "Unable to open magicitems.txt\n");
128         return 1;
129     }
130     BinarySearchTree bst;
131     while (std::getline(magicitems, line)) {
132         bst.insert(line);
133     }
134     magicitems.close();
135     printf("\n");
136
137     //print items in order
138     bst.depth_first_traverse();
139     printf("\n");
140
141     //...
142 }
```

Lines from `magicitems.txt` are read and inserted into the BST. Then `depth_first_search` is called on it.

3.4 BINARY SEARCH TREE SEARCHING

main.cpp

```

138 int main(int argc, char* argv[]) {
139     //...
140
141     //Read lines from magicitems-find-in-bst.txt and search BST
142     std::ifstream file("./magicitems-find-in-bst.txt");
143     if (!file.is_open()) {
144         fprintf(stderr, "Unable to open magicitems-find-in-bst.txt\n");
145         return 1;
146     }
147     int comparisons = 0;
148     int count = 0;
149     while (std::getline(file, line)) {
150         int c = bst.search(line);
151         if (c == -1) {
152             fprintf(stderr, ", Not found after %d comparisons\n", c);
153             return 1;
154         }
155         printf(", Found after %d comparisons\n", c);
156         comparisons += c;
157         count++;
158     }
159     file.close();
160     printf("Average comparisons: %f\n", static_cast<float>(comparisons) /
        ↪ static_cast<float>(count));
161
162     return 0;
163 }

```

Items are read from magicitems-find-in-bst.txt and passed into the BST's search function. Comparisons for each search are returned and averaged.